

## list-test-1-2

May 10, 2025

```
[1]: import tensorflow as tf
print("TensorFlow version:", tf.__version__)
print("GPUs Available:", tf.config.list_physical_devices('GPU'))
```

2025-05-07 16:19:38.892618: I tensorflow/core/platform/cpu\_feature\_guard.cc:182] This TensorFlow binary is optimized to use available CPU instructions in performance-critical operations.

To enable the following instructions: AVX2 FMA, in other operations, rebuild TensorFlow with the appropriate compiler flags.

TensorFlow version: 2.13.0

GPUs Available: [PhysicalDevice(name='/physical\_device:GPU:0', device\_type='GPU'), PhysicalDevice(name='/physical\_device:GPU:1', device\_type='GPU'), PhysicalDevice(name='/physical\_device:GPU:2', device\_type='GPU'), PhysicalDevice(name='/physical\_device:GPU:3', device\_type='GPU')]

```
[2]: import tensorflow as tf

gpus = tf.config.list_physical_devices('GPU')
if gpus:
    try:
        # Select GPU with index 1 (i.e., the second GPU)
        tf.config.set_visible_devices(gpus[3], 'GPU')
        logical_gpus = tf.config.list_logical_devices('GPU')
        print(f"Using GPUs: {logical_gpus}")
    except RuntimeError as e:
        # Visible devices must be set before GPUs have been initialized
        print(e)
else:
    print("No GPUs found.")
```

Using GPUs: [LogicalDevice(name='/device:GPU:0', device\_type='GPU')]

2025-05-07 16:19:41.915259: I

tensorflow/core/common\_runtime/gpu/gpu\_device.cc:1639] Created device

/job:localhost/replica:0/task:0/device:GPU:0 with 30942 MB memory: -> device:

3, name: Tesla V100-DGXS-32GB, pci bus id: 0000:0f:00.0, compute capability: 7.0

[ ]:

[ ]:

```
[3]: import os
import cv2
import numpy as np
import math
import tensorflow as tf
import matplotlib.pyplot as plt
from tensorflow.keras.utils import Sequence
from tensorflow.keras.layers import Input, Conv2D, MaxPooling2D,
↳Conv2DTranspose, concatenate, Dropout
from tensorflow.keras.models import Model
from tensorflow.keras.callbacks import ReduceLROnPlateau, EarlyStopping

# ==== 1. Paths ====
image_dir = '/raid/lits_dataset/train_images/train_images'
mask_dir = '/raid/lits_dataset/train_masks/train_masks'

# ==== 2. Load all valid image-mask pairs ====
all_filenames = [
    f for f in os.listdir(image_dir)
    if f.lower().endswith(('.png', '.jpg', '.jpeg')) and os.path.exists(os.path.
↳join(mask_dir, f))
]

valid_filenames = [
    f for f in all_filenames
    if cv2.imread(os.path.join(mask_dir, f), cv2.IMREAD_GRAYSCALE) is not None
]

# Optional: Print number of positive masks
tumor_count = 0
for f in valid_filenames:
    mask = cv2.imread(os.path.join(mask_dir, f), cv2.IMREAD_GRAYSCALE)
    if np.any(mask > 127):
        tumor_count += 1
print(f"{tumor_count}/{len(valid_filenames)} images contain tumors")

# ==== 3. Train/Val Split ====
np.random.seed(42)
indices = np.arange(len(valid_filenames))
np.random.shuffle(indices)
split_idx = int(0.8 * len(indices))
train_files = [valid_filenames[i] for i in indices[:split_idx]]
val_files = [valid_filenames[i] for i in indices[split_idx:]]
```

0/58638 images contain tumors

```
[4]: # If in a notebook, ensure inline plotting
%matplotlib inline

import os
import cv2
import numpy as np
import matplotlib.pyplot as plt
from random import sample

def overlay_mask_on_image(img, mask, color=(0, 0, 255), alpha=0.5):
    """
    Overlay a single-channel mask onto a BGR image in red.

    Args:
        img : H×W×3 uint8 BGR image.
        mask : H×W uint8 mask (0-255).
        color : BGR tuple for the mask overlay (now red = (0,0,255)).
        alpha : float opacity of the mask overlay [0.0-1.0].
    Returns:
        overlaid : H×W×3 uint8 image.
    """
    mask_color = np.zeros_like(img)
    mask_color[mask > 0] = color
    return cv2.addWeighted(mask_color, alpha, img, 0.8, 0)

# 1. Configure these to point to your folders:
image_dir = '/raid/lits_dataset/train_images/train_images'
mask_dir = '/raid/lits_dataset/train_masks/train_masks'
image_size = (256, 256)

# 2. Find all filenames present in both folders:
all_filenames = [
    f for f in os.listdir(image_dir)
    if f.lower().endswith(('.png', '.jpg', '.jpeg'))
    and os.path.exists(os.path.join(mask_dir, f))
]

print(f"Found {len(all_filenames)} image-mask pairs.")

# 3. Visualize a random subset (up to 5)
n_samples = min(5, len(all_filenames))
for fname in sample(all_filenames, n_samples):
    # Load and resize
    img = cv2.imread(os.path.join(image_dir, fname))
```

```

img = cv2.resize(img, image_size)
mask = cv2.imread(os.path.join(mask_dir, fname), cv2.IMREAD_GRAYSCALE)
mask = cv2.resize(mask, image_size)

# Create overlay in red
overlay = overlay_mask_on_image(img, mask, color=(0, 0, 255), alpha=0.5)

# Convert BGR→RGB for matplotlib
img_rgb = cv2.cvtColor(img, cv2.COLOR_BGR2RGB)
overlay_rgb = cv2.cvtColor(overlay, cv2.COLOR_BGR2RGB)

# Plot original and overlay
plt.figure(figsize=(10, 5))


plt.subplot(1, 2, 1)
plt.title('Original Image')
plt.imshow(img_rgb)
plt.axis('off')

plt.subplot(1, 2, 2)
plt.title('Overlay (Red Mask)')
plt.imshow(overlay_rgb)
plt.axis('off')


plt.show()

```


Found 58638 image-mask pairs.




Original Image




Overlay (Red Mask)



Original Image




Overlay (Red Mask)







Original Image




Overlay (Red Mask)



Original Image



Overlay (Red Mask)



```
[23]: import os
import cv2
import numpy as np
from tensorflow.keras.utils import Sequence
from sklearn.model_selection import train_test_split

# -----
# Your existing DataGenerator class
# -----
class SegmentationDataGenerator(Sequence):
```

```

def __init__(self, image_dir, mask_dir, file_list, batch_size=16,
              image_size=(256,256), shuffle=True, augment_fn=None):
    """
    Args:
        image_dir    : path to folder with input images
        mask_dir     : path to folder with corresponding masks
        file_list    : list of filenames (basename) to include
        batch_size   : number of samples per batch
        image_size   : (height, width) to resize to
        shuffle      : whether to shuffle at epoch end
        augment_fn   : optional function(img, mask) -> (img, mask)
    """
    self.image_dir = image_dir
    self.mask_dir  = mask_dir
    self.file_list = list(file_list)
    self.batch_size = batch_size
    self.image_size = image_size
    self.shuffle    = shuffle
    self.augment_fn = augment_fn
    self.on_epoch_end()

def __len__(self):
    # number of batches per epoch
    return int(np.ceil(len(self.file_list) / self.batch_size))

def __getitem__(self, idx):
    # build batch indexes
    batch_files = self.file_list[idx * self.batch_size : (idx+1) * self.
↪batch_size]

    imgs, masks = [], []
    for fname in batch_files:
        # load & resize
        img = cv2.imread(os.path.join(self.image_dir, fname), cv2.
↪IMREAD_COLOR)
        mask = cv2.imread(os.path.join(self.mask_dir,  fname), cv2.
↪IMREAD_GRAYSCALE)
        img = cv2.resize(img, self.image_size)
        mask = cv2.resize(mask, self.image_size)

        # normalize image
        img = img.astype(np.float32) / 255.0
        # binarize mask
        mask = (mask > 0).astype(np.float32)

        # optional augmentation
        if self.augment_fn is not None:

```

```

        img, mask = self.augment_fn(img, mask)

        imgs.append(img)
        masks.append(mask[..., np.newaxis]) # add channel dim

    X = np.stack(imgs, axis=0)
    y = np.stack(masks, axis=0)
    return X, y

def on_epoch_end(self):
    if self.shuffle:
        np.random.shuffle(self.file_list)

# -----
# 1) Gather and split your filenames
# -----
image_dir = '/raid/lits_dataset/train_images/train_images'
mask_dir = '/raid/lits_dataset/train_masks/train_masks'

# get all matching image-mask names
all_filenames = [
    f for f in os.listdir(image_dir)
    if f.lower().endswith(('.png', '.jpg', '.jpeg'))
    and os.path.exists(os.path.join(mask_dir, f))
]

# split into train/val
train_files, val_files = train_test_split(
    all_filenames,
    test_size=0.2,
    random_state=42,
    shuffle=True
)

print(f"Train samples: {len(train_files)}, Validation samples: {len(val_files)}")

# -----
# 2) (Optional) Define a simple augmentation function
# -----
def simple_augment(img, mask):
    # example: random horizontal flip
    if np.random.rand() < 0.5:
        img = np.fliplr(img)
        mask = np.fliplr(mask)

```



```

    return img, mask

# -----
# 3) Instantiate your train and validation generators
# -----
train_gen = SegmentationDataGenerator(
    image_dir=image_dir,
    mask_dir=mask_dir,
    file_list=train_files,
    batch_size=16,
    image_size=(256,256),
    shuffle=True,
    augment_fn=simple_augment
)

val_gen = SegmentationDataGenerator(
    image_dir=image_dir,
    mask_dir=mask_dir,
    file_list=val_files,
    batch_size=16,
    image_size=(256,256),
    shuffle=False,          # usually no shuffle for validation
    augment_fn=None         # no augmentation on validation
)

```

Train samples: 46910, Validation samples: 11728

[ ]:

```

[19]: import numpy as np
import matplotlib.pyplot as plt

def overlay_mask_on_image(img, mask, color=(1, 0, 0), alpha=0.5):
    overlay = img.copy()
    for c in range(3):
        overlay[..., c] = np.where(
            mask > 0,
            img[..., c] * (1 - alpha) + alpha * color[c],
            img[..., c]
        )
    return overlay

X_val, y_val = val_gen[0]
n = min(4, X_val.shape[0])
plt.figure(figsize=(12, 3 * n))

```

```

for i in range(n):
    img = X_val[i]
    mask = y_val[i, ..., 0] # binary 0/1

    overlay = overlay_mask_on_image(img, mask)


    # Original
    ax = plt.subplot(n, 3, 3*i + 1)
    ax.imshow(img)
    ax.set_title("Original")
    ax.axis('off')

    # Ground truth (white)
    ax = plt.subplot(n, 3, 3*i + 2)
    #-- choose one--#
    ax.imshow(mask, cmap='gray', vmin=0, vmax=1, interpolation='nearest')
    # OR:
    # gt = np.zeros((*mask.shape, 3), dtype=np.float32)
    # gt[mask>0] = 1.0
    # ax.imshow(gt)
    ax.set_title("Ground Truth (white)")
    ax.axis('off')

    # Overlay
    ax = plt.subplot(n, 3, 3*i + 3)
    ax.imshow(overlay)
    ax.set_title("Overlay")
    ax.axis('off')

plt.tight_layout()
plt.show()

```



```
[ ]: from tensorflow.keras.utils import Sequence
import os, cv2
import numpy as np

class SegmentationDataGenerator(Sequence):
    def __init__(self, image_dir, mask_dir, batch_size, image_size):
        self.image_dir = image_dir
        self.mask_dir = mask_dir
```

```

        self.batch_size = batch_size
        self.image_size = image_size

        # automatically grab all filenames
        self.filenames = sorted(os.listdir(self.image_dir))
        # optional: filter to known extensions
        self.filenames = [f for f in self.filenames if f.lower().endswith(('.
↪png', '.jpg', '.jpeg'))]

        assert len(self.filenames) > 0, "No images found in " + image_dir

    def __len__(self):
        return int(np.ceil(len(self.filenames) / self.batch_size))

    def __getitem__(self, idx):
        batch_files = self.filenames[idx * self.batch_size : (idx + 1) * self.
↪batch_size]
        images = []
        masks = []

        # DEBUG: print what we're about to load
        print("Loading batch:", batch_files)

        for fname in batch_files:
            img_path = os.path.join(self.image_dir, fname)
            mask_path = os.path.join(self.mask_dir, fname)
            # Now both img_path and mask_path are guaranteed strings
            img = cv2.imread(img_path)
            mask = cv2.imread(mask_path, cv2.IMREAD_GRAYSCALE)
            img = cv2.resize(img, self.image_size)
            mask = cv2.resize(mask, self.image_size)
            images.append(img)
            masks.append(mask[..., np.newaxis])

        return np.array(images), np.array(masks)

```

```

[16]: import cv2
import matplotlib.pyplot as plt

# helper to overlay mask onto image
def overlay_mask_on_image(img, mask, color=(1, 0, 0), alpha=0.5):
    """
    Overlay a single-channel mask onto an RGB image.

    Args:
        img    : H×W×3 float image in [0,1] RGB
        mask   : H×W binary mask (0 or 1)

```

```

        color : RGB tuple for the mask overlay (default red)
        alpha : float opacity of the mask overlay [0.0-1.0]
Returns:
    overlaid : H×W×3 float image
"""
overlay = img.copy()
for c in range(3):
    overlay[..., c] = np.where(mask > 0,
                                img[..., c] * (1 - alpha) + alpha * color[c],
                                img[..., c])

return overlay

# 1) Pull a batch from val_gen
X_val, y_val = val_gen[0] # shapes: (batch, H, W, 3) and (batch, H, W, 1)

# 2) Visualize the first 4 samples
n = min(4, X_val.shape[0])
plt.figure(figsize=(12, 3 * n))

for i in range(n):
    img = X_val[i] # float RGB in [0,1]
    mask = y_val[i, ..., 0] # binary mask 0/1

    overlay = overlay_mask_on_image(img, mask)

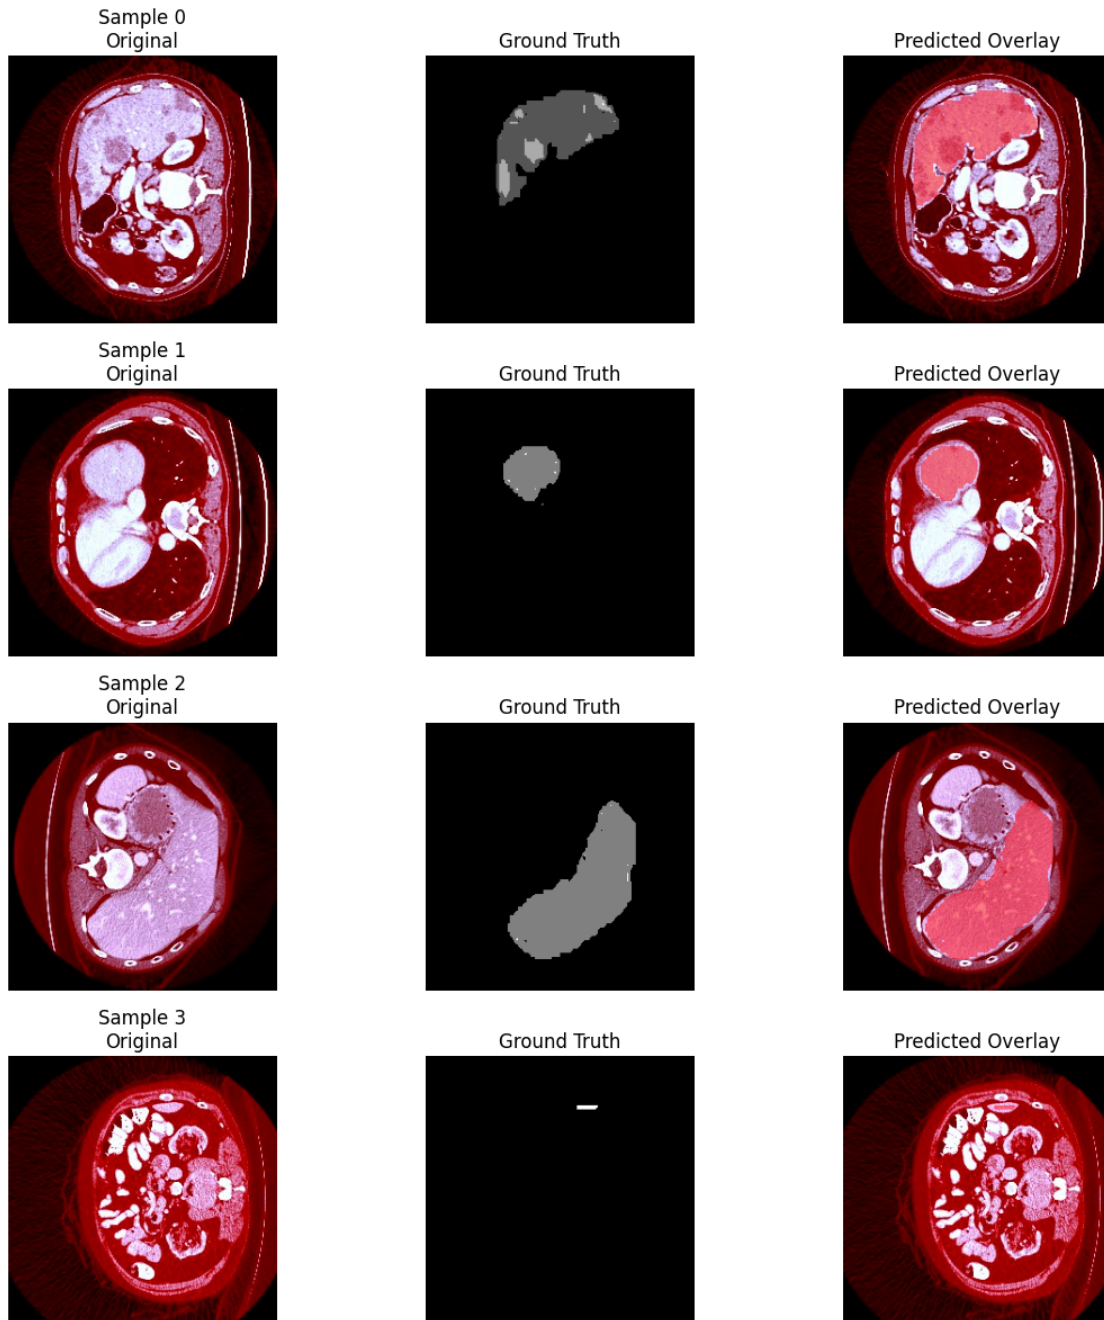
    # Original
    ax = plt.subplot(n, 3, 3*i + 1)
    ax.imshow(img)
    ax.set_title(f"Sample {i}\nOriginal")
    ax.axis('off')

    # Ground truth mask
    ax = plt.subplot(n, 3, 3*i + 2)
    ax.imshow(mask, cmap='gray')
    ax.set_title("Ground Truth")
    ax.axis('off')

    # Overlay
    ax = plt.subplot(n, 3, 3*i + 3)
    ax.imshow(overlay)
    ax.set_title("Predicted Overlay")
    ax.axis('off')

plt.tight_layout()
plt.show()

```



```
[22]: import cv2
import matplotlib.pyplot as plt

# helper to overlay mask onto image
def overlay_mask_on_image(img, mask, color=(1, 0, 0), alpha=0.5):
    """
    Overlay a single-channel mask onto an RGB image.
```



```

Args:
    img      : H×W×3 float image in [0,1] RGB
    mask     : H×W binary mask (0 or 1)
    color    : RGB tuple for the mask overlay (default red)
    alpha    : float opacity of the mask overlay [0.0-1.0]
Returns:
    overlaid : H×W×3 float image
    """
    overlay = img.copy()
    for c in range(3):
        overlay[...] , c] = np.where(mask > 0,
                                     img[...] , c] * (1 - alpha) + alpha * color[c],
                                     img[...] , c])

    return overlay

# 1) Pull a batch from val_gen
X_val, y_val = val_gen[0]    # shapes: (batch, H, W, 3) and (batch, H, W, 1)

# 2) Visualize the first 4 samples
n = min(4, X_val.shape[0])
plt.figure(figsize=(12, 3 * n))

for i in range(n):
    img = X_val[i]            # float RGB in [0,1]
    mask = y_val[i, ..., 0]   # binary mask 0/1

    overlay = overlay_mask_on_image(img, mask)

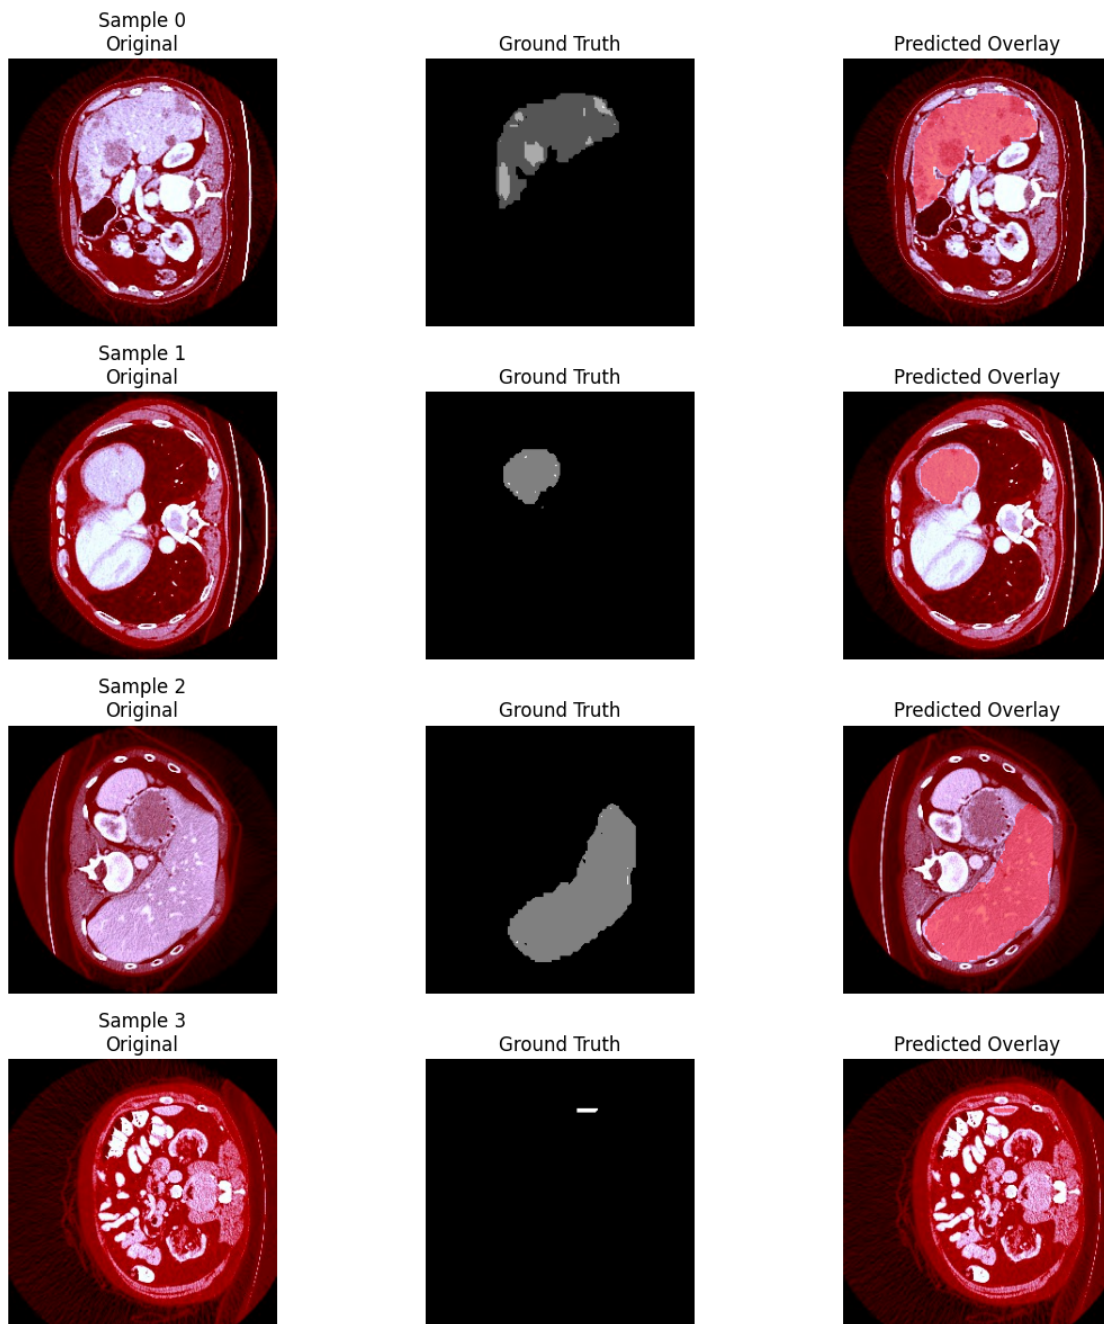
    # Original
    ax = plt.subplot(n, 3, 3*i + 1)
    ax.imshow(img)
    ax.set_title(f"Sample {i}\nOriginal")
    ax.axis('off')

    # Ground truth mask
    ax = plt.subplot(n, 3, 3*i + 2)
    ax.imshow(mask, cmap='gray')
    ax.set_title("Ground Truth")
    ax.axis('off')

    # Overlay
    ax = plt.subplot(n, 3, 3*i + 3)
    ax.imshow(overlay)
    ax.set_title("Predicted Overlay")
    ax.axis('off')

```

```
plt.tight_layout()
plt.show()
```



[ ]:

[ ]: